



Howto set up SGE for CUDA devices?

Asked 13 years, 1 month ago Modified 11 years, 6 months ago Viewed 12k times



5



I'm currently facing the problem of integrating GPU-Servers into an existing SGE environment. Using google I found some examples of Clusters where this has been set up but no information on how this had been done.

Is there some form of howto or tutorial on this anywhere? It doesn't have to be ultra verbose but it should contain enough information to get a "cuda queue" up and running...

Thanks in advance...

Edit: To set up a load sensor about how many GPUs in a node are free, I've done the following:

- set the compute mode of the GPUs to exclusive
- set the GPUs to persistent mode
- add the following script to the cluster configuration as load sensor (and set it so 1 sec.)

```
#!/bin/sh

hostname=`uname -n`

while [ 1 ]; do
    read input
    result=$?
    if [ $result != 0 ]; then
        exit 1
    fi
    if [ "$input" == "quit" ]; then
        exit 0
    fi

    smitool=`which nvidia-smi`
    result=$?
    if [ $result != 0 ]; then
        gpusav=0
        gpus=0
    else
        gpustotal=`nvidia-smi -L|wc -l`
        gpusused=`nvidia-smi |grep "Process name" -A 6|grep -v +-|grep -v \|=|grep -v
Usage|grep -v "No running"|wc -l`
        gpusavail=`echo $gpustotal-$gpusused|bc`
    fi

    echo begin
    echo "$hostname:gpu:$gpusavail"
    echo end
done
```

exit 0

Note: This obviously works only for NVIDIA GPUs

cuda


gridengine

Share Improve this question

edited Oct 18, 2011 at 9:08

asked Oct 17, 2011 at 8:24

Follow

**luxifer**
luxife 177 1 3 12

4 Answers

Sorted by: Highest score (default)

▲

5

▼

🔖

✓

🔄

The strategy is actually fairly simple.

Using `qconf -mc` you can create a complex resource called `gpu` (or whatever you wish to name it). The resource definition should look something like:

#name	shortcut	type	relop	requestable	consumable	default
urgency						
#-----						
gpu	gpu	INT	<=	YES	YES	0 0

Then you should edit your exec host definitions with `qconf -me` to set the number of GPUs on exec hosts that have them:

```
hostname      node001
load_scaling  NONE
complex_values      gpu=2
user_lists      NONE
xuser_lists     NONE
projects        NONE
xprojects        NONE
usage_scaling    NONE
report_variables NONE
```

Now that you've set up your exec hosts, you can request gpu resources when submitting jobs. eg: `qsub -l gpu=1` and gridengine will keep track of how many GPUs are available.

If you have more than one job running per node that uses a GPU you may want to place your GPUs in to exclusive mode. You can do this with the `nvidia-smi` utility.

Share Improve this answer Follow

answered Oct 17, 2011 at 17:02

**Kamil Kisiel**
12.4k 8 50 70

-
- 1 Thank you! That definitely helped me get it working so I gladly accept your answer... There's one downside to this solution though... it depends on what the users say they would use. If you can point me to some elegant solution where the gpu resource is actually monitored and the `-l` parameter is only used to schedule the job instead of also being used to "calculate" how many GPUs are left, that would be great :) – [luxifer](#) Oct 18, 2011 at 7:08
-
- nevermind... I've been able to create a load sensor for this :) – [luxifer](#) Oct 18, 2011 at 9:02
-
- Care to share your load sensor solution? Is there some output of `nvidia-smi` you can monitor to see how many GPUs are actually in use? or? – [Kamil Kisiel](#) Oct 18, 2011 at 20:21
-
- I've already edited my original question to share that solution... the essence of it is: it only works if you put your GPUs in `EXCLUSIVE_THREAD` mode... then you'll get the number of *used* GPUs with the following command: `nvidia-smi |grep "Process name" -A 9|grep -v +-|grep -v \|=|grep -v Usage|grep -v "No running"|wc -l` – [luxifer](#) Oct 19, 2011 at 6:26
-
- 1 You can set the value of "requestable" to be FORCED instead of YES. Then only jobs that specify a value for gpu will be considered to run in the queue. Make sure you set the default to be NONE instead of 0. Also you should just not assign the gpu complex to the non-gpu queue at all. – [Kamil Kisiel](#) Jan 11, 2013 at 23:39
-



5

Open Grid Engine added GPU load sensor support in the 2011.11 release without the need for `nvidia-smi`. The output of the `nvidia-smi` application may (and does) change between driver releases, so the other approach is not recommended.



If you have the GE2011.11 source tree, look for: `dist/gpu/gpu_sensor.c`



To compile the load sensor (need to have the CUDA toolkit on the system):



```
% cc gpu_sensor.c -Invidia-ml
```

And if you just want to see the status reported by the load sensor interactively, compile with:

```
-DSTANDALONE
```

To use the load sensor in a Grid Engine cluster, you will just need to follow the standard load sensor setup procedure:

<http://gridscheduler.sourceforge.net/howto/loadsensor.html>

Sources:

1. <http://marc.info/?l=npaci-rocks-discussion&m=132872224919575&w=2>

Share Improve this answer Follow

answered Feb 14, 2012 at 0:21



[Robert Alexander](#)

51 1 2



4



When you have multiple GPUs and you want your jobs to request a GPU but the Grid Engine scheduler should handle and select a *free* GPUs you can configure a RSMAP (resource map) complex (instead of a INT). This allows you to specify the amount as well as the names of the GPUs on a specific host in the host configuration. You can also set it up as a HOST consumable, so that independent of the slots your request, the amount of GPU devices requested with `-l cuda=2` is for each host 2 (even if the parallel job got i.e. 8 slots on different hosts).

```
qconf -mc
#name          shortcut  type      relop    requestable consumable default
urgency
#-----
-----
      gpu          gpu      RSMAP     <=      YES       HOST      0
0
```

In the execution host configuration you can initialize your resources with ids/names (here simply GPU1 and GPU2).

```
qconf -me yourhost
hostname          yourhost
load_scaling      NONE
complex_values    gpu=2(GPU1 GPU2)
```

Then when requesting `-l gpu=1` the Univa Grid Engine scheduler will select GPU2 if GPU1 is already used by a different job. You can see the actual selection in the `qstat -j` output. The job gets the selected GPU by reading out the `$SGE_HGR_gpu` environment variable, which contains in this case the chose id/name "GPU2". This can be used for accessing the right GPU without having collisions.

If you have a multi-socket host you can even attach a GPU directly to some CPU cores near the GPU (near the PCIe bus) in order to speed up communication between GPU and CPUs. This is possible by attaching a topology mask in the execution host configuration.

```
qconf -me yourhost
hostname          yourhost
load_scaling      NONE
complex_values    gpu=2(GPU1:SCCCCScccc GPU2:SccccSCCCC)
```

Now when the UGE scheduler selects GPU2 it automatically binds the job to all 4 cores (C) of the second socket (S) so that the job is not allowed to run on the first socket. This does not even require the `-binding qsub` param.

More configuration examples you can find on www.gridengine.eu.

Note, that all these features are only available in Univa Grid Engine (8.1.0/8.1.3 and higher), and not in SGE 6.2u5 and other Grid Engine version (like OGE, Sun of Grid Engine etc.). You can try it out by downloading the 48-core limited free version from univa.com.

Share Improve this answer Follow

edited Feb 5, 2013 at 16:00

answered Feb 5, 2013 at 7:10



Daniel

41 2



0



For SGE 2011.11 that comes with ROCKS 6.1 I found that setting the complex consumable to:

#name	shortcut	type	relop	requestable	consumable	default
urgency						
#-----						

gpu	gpu	INT	<=	YES	JOB	0
0						

This allowed me to set the number of GPUs per node and when I submitted a job the number of GPUs requested was not dependent on the SMP/SLOT count. I can then use 8 CPUs and 4 GPUs per job and not cause problems with other jobs leaking in. I still had to set the consumables for the nodes as above.

This is not as nice of a solution as some of the others but I found that the RSMAP option was not available in SGE 2011.11. I would like to eventually get this kind of configuration as I could then set which GPUs get used.

Hope this helps someone save a few hours of configuration.

Share Improve this answer Follow

edited May 9, 2013 at 20:11

answered May 9, 2013 at 19:00



89c3b1b8-b1ae-11e6-b842-48d705

1,475 1 13 22



Thomas J Mustard

1